Qpack2 – a Matlab tool for performing atmospheric retrievals of OEM type

Patrick Eriksson
(patrick.eriksson@chalmers.se)

July 20, 2016

1   Scope

Qpack2 is a retrieval system, implemented in Matlab, for performing inversions of atmospheric observations inside the framework of "optimal estimation" (OEM). With good knowledge of the background systems a wide range of retrievals can be handled, but the primary aim is to provide relatively straightforward retrievals of vertical profiles of atmospheric quantities. That is, the atmosphere is treated to be of "1D type" during the retrieval. Accordingly, a main application areas of Qpack2 is ground-based observations.

The package should be general and flexible for covered measurements. For example, there is basically no limitations regarding observation geometry. Each measurement can consist of several spectra, obtained by some scanning procedure or different instruments. See Sec. ?? for a discussion of possibilities and limitations when it comes to handle more complex measurements. Further, Qpack2 has been developed to be suitable for operational inversions, as long as not most extreme calculation speed is required. Some important aspects here are that batch calculations are allowed and atmospheric a priori profiles can automatically be extracted from climatology data.

2   Background and introduction

Qpack2 is part of the Atmlab package of Matlab functions. In fact, it is largely a merge of some of the systems in Atmlab. The computational engine, i.e. the forward model, of Qpack2 is ARTS-2 (*Eriksson et al.*, 2011). ARTS is a C++ program, downloaded separately (i.e. not part of Atmlab). The communication with ARTS-2 is made through a systen denoted as Qarts. OEM inversions are performed by the function `oem.m`. Climatology data are stored and interpolated through a data format called "atmdata".

This can be seen as a direct successor of Qpack (*Eriksson et al.*, 2005), despite all code is written from scratch. For both Qpack versions the computations are controlled by a set of settings fields packed into a structure denoted as Q (though Qpack2 has also a structure O). Qpack was built around ARTS-1, this new Qpack version is mainly a consequence of that ARTS-2 is the maintained version of ARTS. (ARTS-2 is below just denoted as ARTS.) This gave also an opportunity to make a more stringent implementation and to add some features. This should hopefully make it easier to maintain and extend Qpack2. Data formats are now more clearly defined and the ambition level around documentation is higher (but still modest). A main improvement compared to Qpack is that extraction of a priori data from climatology data is now an integrated part. On the other hand,

the feature of classifying errors into different categories (0-3 in Qpack) and the associated plotting features were removed as these features are difficult to implement in a generic way.

Quality checks are not handled by Qpack2. For example, there is no default check that the frequencies of the measurements are as expected. This for efficiency reasons and the fact that such checks are hard to make in a totally general manner. Checks are left for accompanying dedicated functions, such as `qp2_check_f`.

3   Software and installation

The needed software packages are ARTS and Atmlab. The simplest is to obtain and update these packages through svn. Download instructions for ARTS are found at `www.sat.ltu.se/arts/getarts`. The installation details are here not repeated. A plain installation of ARTS suffices.

Download options for Atmlab are found at `www.sat.ltu.se/arts/tools`. The file CONFIGURE (found in the top folder) gives instructions for how to get started with Atmlab. Qarts2 demands that several "atmlab settings" are activated. This is handled by the function `atmlab`. A description of all atmlab settings is obtained by:

```
>> help atmlab
```

At least the following atmlab settings are used: `ARTS_PATH`, `ARTS_INCLUDES`, `VERBOSITY`, `FMODEL_VERBOSITY`, `STRICT_ASSERT` and `WORK_AREA`. Some example settings:

```
atmlab( 'ARTS_PATH',         fullfile(homedir,'ARTS/arts/src/arts') );
atmlab( 'ARTS_INCLUDES',     fullfile(homedir,'ARTS/arts/includes') );
atmlab( 'FMODEL_VERBOSITY',  0                                      );
atmlab( 'VERBOSITY',         1                                      );
atmlab( 'STRICT_ASSERT',     true                                   );
atmlab( 'WORK_AREA',         '/tmp'                                 );
```

The standard choice is to place these calls of `atmlab` in `atmlab_conf.m`, as described in CONFIGURE.

4   Documentation

The overall documentation is found in this document. The detailed information is mainly stored in the implementation files. Information on individual functions is obtained by the standard Matlab help command. For example:

```
>> help oem
```

Some of the used data structures are documented matching the requirements of `qinfo.m`. For example, to list all fields and the documentation text of Qarts' `Q`:

```
>> qinfo(@qarts)
```

The documentation for a specific field is obtained as

```
>> qinfo(@qarts,'F_BACKEND')
```

Wildcards are allowed when defining fields:

```
>> qinfo(@qarts,'ABS_*')
```

Details around development and bug fixes are described in the ChangeLog file of Atmlab (`atmlab/ChangeLog`).

5 Overview

5.1 Practicalities

A retrieval by Qpack2 has three main steps:

```
% 1: Define forward model and retrieval settings
[Q,O] = my_q_fun;

% 2: Import the measurement data to be inverted
Y = my_y_fun;

% 3: Perform the inversion
L2 = qpack2( Q, O, Y );
```

The variables `Q`, `O`, `Y` and `L2` are all structures, and the fields of these structures are described in Section **??**.

The main part of the settings defined in step 1 are found in `Q`, that match directly the `Q` of the Qarts interface to ARTS. All forward model variables are part of `Q`, and retrieval variables such as grids and a priori data are also defined in `Q`. The exceptions are settings directly associated with OEM, that are stored in `O`. This division reflects the fact that the OEM operations are made by a stand-alone function. This function, `oem`, is implemented in a general manner, to be applicable for any retrieval case as long as an interface to a "forward model" is provided.

The measurement data structure `Y` can be an array. That is, a series of observations can be inverted in a single call of `qpack2`, on the condition that a `Q` and `O` are valid for all observations. This implies e.g. that basic sensor characteristics such as backend frequencies must be common for all the data cases in `Y`. On the other hand, both observation position and direction can differ between the observations, and the data from e.g. an aircraft flight could be inverted in a single run.

The usage of the functions `my_q_fun` and `my_y_fun` in steps 1 and 2 is just a suggestion for how to organise the definition phase of the run, but it is probably a good idea to separate the inversion settings and the task of importing measurement data.

When `Q`, `O` and `Y` are created, it is just to call `qpack2`, and everything should work automatically. The input of `qpack2` is fixed (throughout `Q`, `O` and `Y`). The output is normally the result of the retrieval, packed as a "level 2" data structure (`L2`). The function can also provide simulated measurements. This option is triggered by setting the spectrum field to be empty (if `Y` is an array, this must be done for all elements):

```
...
Ye.Y = [];
Ysim = qpack2( Q, O, Ye );
```

The output is a copy of `Ye`, with the field `Ysim.Y` set to the result of the forward model run. The simulated data match the a priori state of the retrieval set-up.

Several full examples are found in Atmlab's folder `demos`, such as `qpack2_demo.m`. The examples work with simulated data, created by `qpack2` itself. The `qpack2_demo.m` example file is also appended to this document as Appendix **??**. A related example is `arts_oem_demo`. It does not use Qpack2, but shows how an OEM inversion is performed with ARTS as the forward model, and the calculation steps in this example gives a good view of the interior of Qpack2.

5.2 Units

In most cases SI units are used. For example, frequency off-sets are retrieved in Hz. For retrieval of species profiles several "units" can be used. First of all, both volume mixing ratio and number density retrievals can be made. These options are selected by setting the species unit to `vmr` and `nd`, respectively.

Further, the species unit can be also be set to `rel`. In this case, the retrieved profile is expressed in fractions of the a priori profile. For example, the a priori state corresponds to 1, and $x = 2$ signifies that the species amount is double as high as the a priori. The remaining species unit is `logrel`, that is defined as the natural logarithm of the `rel` unit. The main reason for using this unit is to introduce a constrain of positive species values. For example, if z is a retrieved `logrel` value and v_a is the a priori VMR, the retrieved volume mixing ratio is $v_a \exp(z)$, which always is greater than 0.

The prescribed or selected unit is used for all related variables. For example, a priori covariance matrices must be specified correspondingly. For example, if the selected species unit is `rel`, the 1 standard deviation a priori uncertainty should be in the order of 0.5 (that corresponds to a 50% uncertainty). The same applies to output data. No automatic unit conversions are performed (in contrast to the earlier Qpack version), e.g. the returned a priori profile for `logrel` retrievals is a vector of zeros. The basic idea is to return data that are fully consistent with the assumptions of the retrieval.

Data retrieved using the `rel` and `logrel` units can be converted to volume mixing ratio data by the function `qp2_rel2vmr`. However, this function handles only data directly associated with the species, it does not cover, for example, the Jacobian matrix. See further the help text of the function.

6 Main data structures

6.1 Forward model and retrieval settings

6.1.1 General features

Forward model parameters and most retrieval variables are joined into a single structure, denoted as `Q`. This structure has a fixed set of fields. That is, the structure must always contain exactly this set of fields. However, some fields are used only in particularly circumstances and not all fields must be set. A field of `Q` is flagged as undefined by setting it to `{}`, and e.g. `[]` is taken as an active selection. This applies also to structures `O` and `Y`.

The structure `Q` of Qpack2 is identical to the `Q` of Qarts. Qpack2 makes only slight modifications of `Q` before it is given to Qarts. The Qarts system is primarily an interface to the ARTS forward model, on the same time as retrieval variables can be defined in a relatively general manner. However, Qarts is not a complete retrieval system, it just allows definition of retrieval variables in parallel with the forward model data. The ambition is just to provide the basis for retrieval systems, such as Qpack2 or streamlined software dedicated to a single sensor.

Qarts is documented through the qinfo feature:

```
>> qinfo(@qarts)
```

The call of `qinfo` above provides a description of all defined fields (Sec. **??**) and this information is also found here as Appendix **??**. It is recommended to initialise `Q` as

```
Q = qarts;
```

This ensures that Q contains all defined fields, also ones introduced more lately.

All features of ARTS can be accessed through Qarts. Calculations of standard type can largely be handled by defining the fields of Qarts that have a direct matching ARTS workspace variables (WSV). That is, the name is the same in Qarts and ARTS, beside that Qarts use uppercase for its fields and ARTS lowercase for its variables (e.g. `Z_SURFACE` and `z_surface`, respectively). Qarts has also several fields that allow inclusion of workspace method (WSM) calls. Some of these fields match ARTS agendas and have then a dedicated purpose (e.g. `SURFACE_RTPROP_AGENDA`). There are also fields where you are free to include any set of WSM calls (e.g. `WSMS_BEFORE_RTE`). You can also use ARTS's include feature through the field `INCLUDES`.

Accordingly, Qarts is a relatively clean and open interface to ARTS, and to set-up a calculation demands primarily an understanding of ARTS. Built-in documentation of ARTS WSMs and WSVs is obtained in a terminal as

```
$ arts -d z_surface
```

or from within Matlab as

```
>> arts('-d z_surface');
```

See the documentation of ARTS for more in-depth documentation.

The best introduction to the practical usage of ARTS is obtained through the control file examples found in ARTS's folder `controlfiles`. For an introduction of the usage of ARTS through Qarts, there are several example scripts in Atmlab's folder `demos`, such as `qarts_demo`.

6.1.2 Qpack2 specifics

Most fields of `Q` can be unspecified (`{}`), but Qpack2 handles some fields in special way:

- `Q.RAW_ATMOSPHERE` must be unset (`{}`).

- All existing "atmdata"-fields are considered and are given highest priority. That is, if such a field is set, any setting of the corresponding field will be overwritten. For example, if `Q.T.ATMDATA` is set, any existing values in `T_FIELD` will have no effect.

- The fields `T_FIELD`, `VMR_FIELD` and `Z_FIELD` must be defined directly in `Q`, possibly through the corresponding atmdata fields (see point above). That is, these workspace variables can not be set by an include file.

- In the case of `Z_FIELD` there is another option, this field can be set by applying hydrostatic equilibrium, then activated by setting `Q.HSE.ON` to true. The reference point for hydrostatic equilibrium is specified by `Y.HSE_P` and `Y.HSE_Z` (Sec. **??**).

- `Q.HSE_P` is set to `Y.HSE_P`. Other fields of `Q.HSE` must be set by the user, thus including an active choice for `Q.HSE.ON`.

- `Q.ABSORPTION` must be set. The option 'CalcTable' is not allowed (as it is difficult to ensure generation of good absorption tables in a general manner). You must either use 'OnTheFly' or create an absorption table as a pre-calculation (the 'LoadTable' option).

- `Q.TNOISE_C` must be set.

- The a priori knowledge covariance matrix (Sx) is created by `arts_sx` and the `SX` sub-fields of the retrieval quantities must be set.

6.2 OEM settings

Settings directly associated with OEM are put into the structure O. Also this setting structure is documented through `qinfo` (see also Appendix **??**):

```
>> qinfo(@oem)
```

Many of the fields of O controls the output of the `oem` function. For Qpack2 you do not need to set these fields manually. The output required to provide the specified L2 data (Section **??**) is ensured by initialising O as

```
O = qp2_l2( Q );
```

Some fields of O are also set by Qpack2, and user settings are overwritten. These fields are:

- O.msg

- O.sxnorm

For linear inversion, a single active setting is needed:

```
O.linear = true;
```

6.3 Measurement data

The series of measurements to be treated is packed into a data structure denoted as Y. The data fields of this structure are described in `qp2_y`, following the `qinfo` format:

```
>> qinfo(@qp2_y)
```

The same information is appended to this document in Appendix **??**. An example:

```
>> Y(3)

ans =

          DAY: 3
            F: [800x1 double]
         HOUR: 11
        HSE_P: 10000
        HSE_Z: 1.5677e+04
     LATITUDE: 57.4000
    LONGITUDE: 11.9300
       MINUTE: 57
        MONTH: 3
       SECOND: 58.0000
       TNOISE: 0.0243
            Y: [800x1 double]
         YEAR: 2002
           ZA: 65
    Z_PLATFORM: 5
```

The simplest way to initialise Y is to use `qp2_y` as

```
for i = 1 : n
  Y(i) = qp2_y;
  load( file{i} )
  Y(i).Y = ...
end
```

6.4   Retrieved data

Retrieved data are returned as a structure array `L2`. The fields of `L2` are not fixed, the set of returned fields depend on settings in `Q`. This makes it impossible to use qinfo here (as for `Q` and `Y`), and the fields are instead described below in this document.

The only mandatory field of `L2` is `converged`. The remaining content of L2 is controlled by the Qarts field `L2_EXTRA` and the `L2` sub-field for the retrieval quantities (e.g. `Q.POLYFIT.L2`). As an example on a complete L2, the first element of the output of `qpack2_demo` is (the polyfit part is truncated):

```
>> L2(1)

ans =

              year: 2008
             month: 2
               day: 25
              hour: {}
            minute: {}
            second: {}
         converged: 1
                dx: 4.4374e-06
              cost: 1.0748
            cost_x: 5.8390e-04
            cost_y: 1.0742
                 f: [1279x1 double]
                 y: [1279x1 double]
                yf: [1279x1 double]
                bl: [1279x1 double]
            p_grid: [45x1 double]
           t_field: [45x1 double]
           z_field: [45x1 double]
     species1_name: 'O3'
        species1_p: [45x1 double]
       species1_xa: [45x1 double]
        species1_x: [45x1 double]
        species1_e: [45x1 double]
       species1_eo: [45x1 double]
       species1_es: [45x1 double]
       species1_mr: [45x1 double]
           ffit_xa: 0
            ffit_x: -3.8104e+03
            ffit_e: 8.7798e+03
           ffit_eo: 8.6418e+03
```

```
        ffit_es: 1.5504e+03
        ffit_mr: 0.9692
     polyfit0_xa: 0
      polyfit0_x: 0.0516
      polyfit0_e: 0.4748
     polyfit0_eo: 0.0967
     polyfit0_es: 0.4649
     polyfit0_mr: 0.7745
     polyfit1_xa: 0
      polyfit1_x: -0.0022
      polyfit1_e: 0.0097
     polyfit1_eo: 0.0094
     polyfit1_es: 0.0023
     polyfit1_mr: 0.9996
        ...
```

A description of the options for `Q.L2_EXTRA` is obtained by:

```
>> help qp2_l2
```

Further comments and description of fields directly associated with specific retrieval quantities are found below (no ambition of completeness here).

6.4.1 Various general fields

`cost`, `cost_x` and `cost_y`:
Cost values of retrieved state and fit to measurement. See further `oem.m`. All obtained by including the string 'cost' in `L2_EXTRA`.

`P.year`, `P.month`, `P.day`, `P.hour`, `P.minute`, and `P.second`:
The time of the measurement. All obtained by including the string 'date' in `L2_EXTRA`.

`tnoise`:
Mean of assumed thermal noise. This is simple average, no weighting with channel widths is applied. This value has no direct physical meaning, but can be used as a mean to e.g. remove the most noisy measurements (without having `Y` at hand). Obtained by including the string 'tnoise' in `L2_EXTRA`.

`f`:
Frequency vector of measurement. Equals `Y.F` if set. Otherwise set `Q.SENSOR_RESPONSE_F` or `Q.F_GRID`, depending on if a sensor is applied or not. Only included together with `y`.

`y`:
Measured spectrum. Obtained by including the string 'y' in `L2_EXTRA`.

`yf`:
Fitted spectrum. That is, the spectrum matching retrieved state vector. Obtained by including the string 'yf' in `L2_EXTRA`.

`bl`:
Retrieved "baseline". That is, retrieved spectrum distortion through e.g. polynomial fitting. A scalar 0 if no baseline retrieval is performed. Obtained by including the string 'bl' in `L2_EXTRA`.

`converged`:
Convergence flag. See further `oem.m`. Always included.

`dx`:
Change in the state vector x between each iteration. See further `oem.m`. Obtained by including the string 'dx' in `L2_EXTRA`.

`J`, `G` and `A`:
The (full) Jacobian, gain and averaging kernel matrix, respectively. Obtained by including the string 'J', 'G' and 'Afull', respectively, in `L2_EXTRA`. Note that the string ' A' only triggers inclusion of partial averaging kernels (see below).

`S`, `So` and `Ss`:
The covariance matrix of total, observation and smoothing error, respectively. Obtained by including the string 'S', 'So' and 'Ss', respectively, in `L2_EXTRA`.

`jq` and `ji`:
Information describing the different retrieval quantities. The structure `jq` is basically matching the user settings. One example:

```
>> L2(1).jq{1}

ans =


        maintag: 'Absorption species'
         subtag: 'O3'
           mode: 'vmr'
     analytical: 1
   perturbation: 0
          grids: {[48x1 double]}
```

The second structure, `ji`, gives the start and end index in the state vector for each retrieval quantity. The data are stored in a manner matching the data type used in ARTS for this information, an ArrayOfIndex. For example, if retrieval quantity two holds position 49 in 96 in the state vector, we have that

```
>> L2(1).ji{2}

ans =

   [49]    [96]
```

These two structures are included automatically when `A`, `S` or any data matching the complete state vector is included in the output, and, hence, the structures are needed match vector and matrix elements to the different retrieval quantities.

6.4.2 Background atmospheric state

`p_grid`, (`lat_grid`, `lon_grid`), `t_field` and `z_field`:
As the ARTS workspace variables with the same name. The data are identical to the ones that are provided to the forward model. (If temperatures are retrieved, this information is used. Otherwise returned data match climatology temperatures.) Obtained by including the string 'ptz' in `L2_EXTRA`.

6.4.3   Retrieved species

The retrieved species are indexed after the order they are specified in `Q` and named as `species1`, `species2` ... A number of fields is returned for each species. The field names below are valid for species 1. If two species are retrieved there is also a field named as `species2_name` etc. The unit of output follows `Q.SPECIES(i).UNIT`, see further Section ??. These set of fields are available:

`species1_name:`
Tag name of species.

`species1_p:`
Vertical/pressure retrieval grid for the species.

`species1_x:`
Retrieved profile. Unit depends on `Q.SPECIES(i).UNIT`.

`species1_xa:`
A priori state for the species. Unit depends on `Q.SPECIES(i).UNIT`. Obtained by including the string 'xa' in `L2_EXTRA`.

`species1_e:`
Total retrieval (observation + smoothing) error for the species. Unit depends on `Q.SPECIES(i).UNIT`. Obtained by including the string 'e' in `L2_EXTRA`.

`species1_eo:`
Observation error for the species. Unit depends on `Q.SPECIES(i).UNIT`. Obtained by including the string 'eo' in `L2_EXTRA`.

`species1_es:`
Smoothing error for the species. Unit depends on `Q.SPECIES(i).UNIT`. Obtained by including the string 'es' in `L2_EXTRA`.

`species1_mr:`
Measurement response for the species. This is the sum of the rows of `species_A` (see below). Obtained by including the string 'mresp' in `L2_EXTRA`.

`species1_A:`
The species specific averaging kernels. That is, this is the diagonal centred sub-matrix covering only the resolution between changes of the species at different altitudes. The resolution between variables of different retrieval quantities is thus not covered. Obtained by including the string 'A' in `L2_EXTRA`.

`species1_vmr0:`
Added automatically if the retrieval unit is `rel` or `logrel`, to be used to convert data to VMR. For example, the retrieved profile in VMR for the `rel` case is `x.*vmr0`.

6.4.4   Baseline fit variables

There are two approaches for performing baseline fits, denoted as `POLYFIT` and `SINEFIT`. The two approaches can be used seperately, or be combined.

A set of fields is added for each polynomial coefficient, and for the first polynomial coefficient they are

`P.polyfit0_x`, `P.polyfit0_xa`, `P.polyfit0_e`, `P.polyfit0_eo`, `P.polyfit0_es` and `P.polyfit0_mr:`

As corresponding fields above for species (`species1_x` etc.).

Fields are added in a similar manner for `SINEFIT`. For the first selected ripple period the fields are

`P.sinefit1_x`, `P.sinefit1_xa`, `P.sinefit1_e`, `P.sinefit1_eo`, `P.sinefit1_es` and
`P.sinefit1_mr`:
These fields have length 2, where the two values correspond to the cosine and sine terms, respectively.

6.4.5 Frequency fit variables

Both a shift and a stretch factor can be retrieved:

`P.fshift_x`, `P.fshift_xa`, `P.fshift_e`, `P.fshift_eo`, `P.fshift_es` and
`P.fshift_mr`:
As corresponding fields above for species (`species1_x` etc.).

`P.fstretch_x`, `P.fstretch_xa`, `P.fstretch_e`, `P.fstretch_eo`, `P.fstretch_es` and
`P.fstretch_mr`:
As corresponding fields above for species (`species1_x` etc.).

7 Data types and file formats

7.1 Qpack2 variables

The data types in Qpack2 (and in Qarts) follow largely ARTS. Regarding basic data types, there are two main things to consider. Firstly, a distinction is made between the logical 0 and 1, and the integers 0 and 1. Accordingly, booleans, sometime also denoted as flags, must be set to `true` or `false`. Secondly, vectors are demanded to be columns.

A notable feature of Qpack2 and Qarts is that data fields, booleans and scalars excluded, can either be set directly or be set to the name of a file (including data of the expected type). Both these settings are possible

```
Q.F_GRID = [ 501.18e9 : 1e6 : 501.58e9 ]';
Q.F_GRID = 'f_odinsmr_ac2a.xml;
```

Note the transpose in the first case (to create a column vector).

The fields `Q.ABS_SPECIES.ATMDATA`, `Q.T.ATMDATA` and `Q.Z_ATMDATA` are intended for data of climatology character. The data format used here is called atmdata. The format is similar to the GriddedField format in ARTS, but has also fields for e.g. units. The atmdata format is based on the general gformat:

```
>> help gformat
```

For a description of the specialities of atmdata:

```
>> help isatmdata
```

For example, the MSIS temperature climatology found in arts-xml-data is stored as GriddedField4 but imported as

```
Q.T_ATMDATA = gf_artsxml( fullfile( arts_xmldata_path, 'climatology', ...
                          'msis90', 'msis90.t.xml' ), 'Temperature', 't_field' );
```

This produces data of atmdata type:

```
>> Q.T_ATMDATA

ans =

          TYPE: 'atmdata'
          NAME: 'Temperature'
        SOURCE: [1x65 char]
           DIM: 4
          DATA: [4-D double]
     DATA_NAME: 'Temperature'
     DATA_UNIT: 'K'
         GRID1: [141x1 double]
    GRID1_NAME: 'Pressure'
    GRID1_UNIT: 'Pa'
         GRID2: [19x1 double]
    GRID2_NAME: 'Latitude'
    GRID2_UNIT: 'deg'
         GRID3: 0
    GRID3_NAME: 'Longitude'
    GRID3_UNIT: 'deg'
         GRID4: [13x1 double]
    GRID4_NAME: 'doy'
    GRID4_UNIT: ''
```

7.2 File formats

The standard format for data input files is "ARTS xml". These files are hardly created by hand, but stored from Matlab (through `xmlStore`) or from ARTS (through `WriteXML`). If ARTS is compiled with support for NetCDF, that format can also be used (`arts_nc_write_`*datatype* and `WriteNetCDF`, respectively).

8 Measurements consisting of multiple spectra

As long as each measurement consists of a single spectrum, Qpack2 should be quite flexible in treating sensor aspects. For some conditions, it also possible to handle measurements consisting of multiple spectra.

8.1 Common sensor characteristics

If the sensor characteristics are common for all spectra involved (not including thermal noise) it is relatively straightforward to apply Qpack. For example, the same sensor is moved vertically, or rotated to measure at different zenith angles, can be treated. Requirements and considerations to handle such multiple-spectra measurements include:

- The fields `Y.Z_PLATFORM` and `Y.ZA` shall give the altitude and zenith angle of each spectrum, and are accordingly (column) vectors instead of scalar values for single-spectra cases.

- `Y.LATITUDE` and `Y.LONGITUDE` are kept as scalar values, the same geographical position is assumed for all spectra.

- `Y.T_NOISE` must be extended where each column describes the noise of the corresponding spectrum. Hence, the noise level can vary between the spectra.

- If a baseline fit is performed, the corresponding `SX`-fields must be extended. These fields are scalar values for single-spectrum observations. With multiple spectra, the covariance matrices can include an expected correlation between the baseline coefficient for different spectra. For example, if the zero-order `POLYFIT` coefficient is expected to have a standard deviation of $2\,\mathrm{K}$ and has a correlation of 0.9 between two spectra, these assumptions are included as

  ```
  Q.POLYFIT.SX0 = 2.0^2 * [ 1 0.9; 0.9 1 ];
  ```

It is also possible to describe the scanning in a 3D atmosphere, that should be of interest if the azimuth angle matters. This option should be of interest for cases including a Zeeman influence or wind retrievals. Additional considerations are then:

- `Q.ATMOSPHERE_DIM` shall be set to 3.

- `Q.LAT_GRID` and `Q.LON_GRID` must be specified.

- Atmospheric fields, such as `Q.T_FIELD`, must be 3D. The standard choice should be to set the fields to be constant in latitude and longitude directions. Using the `ATMDATA` feature with an input 1D atmosphere handles this automatically.

- The field `Y.AA` shall be set, and match `Y.ZA` in length.

- Latitude and longitude retrieval grids, such in `Q.ABS_SPECIES.GRIDS`, must be specified, but must have length one. Any value should work. One example:

  ```
  Q.ABS_SPECIES(1).GRIDS = { Q.P_GRID, mean(Q.LAT_GRID), mean(Q.LON_GRID)};
  ```

  Retrieved atmospheric data will effectively have no latitude and longitude dimension and be of 1D type.

The file `qpack2_wind3d_demo` exemplifies retrievals of this type.

8.2 More complex set-ups

It can be considered to append spectra from different instruments to perform a joint retrieval. In a few cases this can be handled as in the sub-section above, but would normally require a more complex set-up. Another more complex situation is when the sensor characteristics change during e.g. a limb scanning sequence. Both these examples can be handled by ARTS by making use of the `yCalcAppend` workspace method.

Using Qpack together with `yCalcAppend` is not impossible but requires detailed knowledge of how both Qpack and ARTS work and far from all possible measurement scenarios can be treated. Hence, it is best to contact Patrick for advice before start trying to implement a set-up of this type.

References

Eriksson, P., C. Jiménez, and S. A. Buehler, Qpack, a general tool for instrument simulation and retrieval work, *J. Quant. Spectrosc. Radiat. Transfer*, 91, 47-64, 2005.

Eriksson, P., S. A. Buehler and C. P. Davis and C. Emde and O. Lemke, ARTS, the atmospheric radiative transfer simulator, Version 2, *J. Quant. Spectrosc. Radiat. Transfer*, 112, 1551-1558, 2011.

Appendices

A   An example script

```
% QPACK2_DEMO   Demonstration of the Qpack2 retrieval system
%
%    The main features of Qpack2 are demonstrated. The example case is airborne
%    measurements of ozone at 110.8 GHz. Synthetic measurement data are
%    generated internally. See the code and internal comments for details.
%
%    Everything is here put into a single file. For practical retrievals it
%    is probably better to put the definitions of Q (together with O?) in a
%    separate function (i.e. [Q,O] = q_mycase). A function to import
%    measurement data into the "Y format" (see *qp2_y*) is needed. The
%    retrieval result is returned in the L2 format produced by *qp2_l2*.
%
%    This script focuses on giving an introduction, indicating different
%    retrieval units and retrieval of other variables. See also *qpack2_demo2*.
%
% FORMAT   L2 = qpack2_demo
%
% OUT     L2   L2 data output from *qpack2*.

% 2010-05-12   Created by Patrick Eriksson.

function L2 = qpack2_demo

errid = ['atmlab:' mfilename];

%- Qarts settings
%
Q    = q_demo;              % Local file, found below


%- Measurement data
%
Y = y_demo( Q );           % Local file, found below


%- Check that all frequencies are OK
%
if ~qp2_check_f( Q, Y, 1e3 );
  error( errid, ...
            'Some mismatch between Q.F_BACKEND and frequencies of spectra.' );
end


%- OEM variables
%
O = qp2_l2( Q );           % This ensures that OEM returns the variables needed
%                            to fill the L2 structure, as defined in Q
```

```
O.linear = false;
%
if ~O.linear
  O.itermethod = 'GN';
  O.stop_dx    = 0.01;
  O.maxiter    = 5;
end



%- Make inversion
%
L2 = qpack2( Q, O, Y );



%- Plot, if no output argument
%
if ~nargout

  % Profiles
  figure(1),clf
  plot( L2(1).species1_x*1e6, p2z_simple(L2(1).species1_p)/1e3, 'b', ...
        L2(2).species1_x*1e6, p2z_simple(L2(2).species1_p)/1e3, 'r', ...
        L2(1).species1_xa*1e6, p2z_simple(L2(1).species1_p)/1e3, 'k-' );
  xlabel( 'Ozone [VMR]' );
  ylabel( 'Approximate altitude [km]' );
  axis( [ 0 8 10 90 ] )
  legend( 'Retrieval 1', 'Retrieval 2', 'True and a priori' );

  % Spectra
  figure(2),clf
  h=plot( L2(1).f/1e9, L2(1).y, 'k.', L2(2).f/1e9, L2(2).y, 'k.', ...
          L2(1).f/1e9, L2(1).yf, 'b-', L2(2).f/1e9, L2(2).yf, 'r-', ...
          L2(1).f/1e9, L2(1).bl, 'b-.', L2(2).f/1e9, L2(2).bl, 'r-.' );
  xlabel( 'Frequency [GHz]' );
  ylabel( 'Tb [K]' );
  axis( [ min(L2(1).f/1e9) max(L2(1).f/1e9) 0 18 ] )
  legend( h(3:end), 'Fitted 1', 'Fitted 2', 'Baseline 1', 'Baseline 2' );
end

return
%------------------------------------------------------------------------------
%------------------------------------------------------------------------------
%------------------------------------------------------------------------------




function Q = q_demo
```

```
errid = ['atmlab:' mfilename];
%- Atmlab settings
%
arts_xmldata_path = atmlab( 'ARTS_XMLDATA_PATH' );
arts_includes     = atmlab( 'ARTS_INCLUDES' );
if isnan( arts_xmldata_path )
  error( errid,'You need to set ARTS_XMLDATA_PATH to run this exmaple.' );
end
if isnan( arts_includes )
  error( erird,'You need to ARTS_INCLUDES to run this example.' );
end
%
fascod = fullfile( arts_xmldata_path, 'planets', 'Earth', 'Fascod' );


%- Init Q
%
Q = qarts;
%
Q.INCLUDES              = { fullfile( 'ARTS_INCLUDES', 'general.arts' ), ...
                            fullfile( 'ARTS_INCLUDES', 'agendas.arts' ), ...
                            fullfile( 'ARTS_INCLUDES', 'continua.arts' ), ...
                            fullfile( 'ARTS_INCLUDES', 'planet_earth.arts' ) };
Q.ATMOSPHERE_DIM       = 1;
Q.STOKES_DIM           = 1;
Q.J_DO                 = true;
Q.CLOUDBOX_DO          = false;


%= Define agendas
%
% Here we do it by using the predefined agenda templates
%    (found in arts/controlfiles/general/agendas.arts)
% This works only if the pre-defined agenda is names following the pattern:
%    name_of_agenda__(Something)
%
Q.PPATH_AGENDA                = { 'ppath_agenda__FollowSensorLosPath'   };
Q.PPATH_STEP_AGENDA           = { 'ppath_step_agenda__GeometricPath'    };
Q.BLACKBODY_RADIATION_AGENDA = { 'blackbody_radiation_agenda__Planck'  };
Q.IY_SPACE_AGENDA             = { 'iy_space_agenda__CosmicBackground'   };
Q.IY_SURFACE_AGENDA           = { 'iy_surface_agenda__UseSurfaceRtprop' };
Q.IY_MAIN_AGENDA              = { 'iy_main_agenda__Emission'            };


%- Radiative transfer
%
Q.IY_UNIT               = 'RJBT';
Q.YCALC_WSMS            = { 'yCalc' };
```

```
%
Q.PPATH_LMAX            = 250;



%- Surface
%
Q.Z_SURFACE             = 1e3;              % Just a dummy value. A 10 km
                                            % observation altitude is assumed here



%- Absorption
%
Q.ABS_LINES             = fullfile( atmlab_example_data, 'o3line111ghz' );
Q.ABS_LINES_FORMAT      = 'Arts';
%
Q.ABSORPTION            = 'OnTheFly';
Q.ABS_NLS               = [];

%- Pressure grid
%
z_toa                   = 95e3;
%
Q.P_GRID                = z2p_simple( Q.Z_SURFACE-1e3 : 2e3 : z_toa )';



%- Frequency, spectrometer and pencil beam antenna
%
% The hypothetical spectrometer has rectangular response functions
%
Q.F_GRID                = qarts_get( fullfile( atmlab_example_data , ...
                                            'f_grid_111ghz.xml' ) );
%
H                       = qartsSensor;
%
H.SENSOR_NORM           = true;
%
H.BACKEND_DO            = true;
df                      = 0.5e6;
H.F_BACKEND             = ( min(Q.F_GRID)+df : df : max(Q.F_GRID)-df )';
%
B.name                  = 'Spectrometer channel response function';
B.gridnames             = { 'Frequency' };
B.grids                 = { [-df/2 df/2] };
B.dataname              = 'Response';
B.data                  = [1 1];
%
H.BACKEND_CHANNEL_RESPONSE{1} = B;
clear B
```

```
%
Q.SENSOR_DO          = true;
Q.SENSOR_RESPONSE    = H;
%
Q.ANTENNA_DIM        = 1;
Q.MBLOCK_ZA_GRID     = 0;




%- Correlation of thermal noise
%
f              = H.F_BACKEND;
cl             = 1.4 * ( f(2) - f(1) );
cfun           = 'gau';
cco            = 0.05;
%
Q.TNOISE_C     = covmat1d_from_cfun( f, [], cfun, cl, cco );
%
clear H f




%- Define L2 structure (beside retrieval quantities below)
%
Q.L2_EXTRA     = { 'cost', 'dx', 'xa', 'y', 'yf', 'bl', 'ptz', ...
                   'mresp', 'A', 'e', 'eo', 'es', 'date' };




%- Temperature
%
Q.T.RETRIEVE   = false;
Q.T.ATMDATA    = gf_artsxml( fullfile( arts_xmldata_path, 'climatology', ...
                   'msis90', 'msis90.t.xml' ), 'Temperature', 't_field' );

%- Determine altitudes through HSE
%
Q.HSE.ON       = true;
Q.HSE.P        = Q.P_GRID(1);
Q.HSE.ACCURACY = 0.1;



%- Species

% Ozone, only species is retrieved here
Q.ABS_SPECIES(1).TAG      = { 'O3' };
Q.ABS_SPECIES(1).RETRIEVE = true;
Q.ABS_SPECIES(1).L2       = true;
Q.ABS_SPECIES(1).GRIDS    = { Q.P_GRID, [], [] };
Q.ABS_SPECIES(1).ATMDATA  = gf_artsxml( fullfile( fascod, ...
     'midlatitude-winter', 'midlatitude-winter.O3.xml' ), 'O3', 'vmr_field' );
```

```
%
% If you don't apply a min value (by MINMAX), you could need to active this:
%Q.VMR_NEGATIVE_OK = true;
%
% For demonstration, setting for several units are provided:
switch 1
 case 1  % Constant VMR
  Q.ABS_SPECIES(1).UNIT      = 'vmr';
  Q.ABS_SPECIES(1).SX        = ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, 1.5e-6, ...
                                          'lin', 0.2, 0.00, @log10 ) + ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, 0.3e-6, ...
                                          'lin', 0.5, 0.00, @log10 );
   Q.ABS_SPECIES(1).MINMAX  = 1e-12;
 case 2 % Constant rel
  Q.ABS_SPECIES(1).UNIT      = 'rel';
  Q.ABS_SPECIES(1).SX        = ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, 0.5, ...
                                          'lin', 0.2, 0.00, @log10 ) + ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, 0.1, ...
                                          'lin', 0.5, 0.00, @log10 );
  Q.ABS_SPECIES(1).MINMAX  = 1e-6;
 case 3 % Mimic case 2 in vmr
  Q.ABS_SPECIES(1).UNIT      = 'vmr';
  Q.ABS_SPECIES(1).SX        = ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, ...
                                   [ Q.ABS_SPECIES(1).ATMDATA.GRID1,...
                                 0.5 * Q.ABS_SPECIES(1).ATMDATA.DATA ],...
                                          'lin', 0.2, 0.00, @log10 ) + ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, ...
                                   [ Q.ABS_SPECIES(1).ATMDATA.GRID1, ...
                                 0.1 * Q.ABS_SPECIES(1).ATMDATA.DATA ],...
                                          'lin', 0.5, 0.00, @log10 );
  Q.ABS_SPECIES(1).MINMAX  = 1e-12;
 case 4 % Constant logrel
  Q.ABS_SPECIES(1).UNIT      = 'logrel';
  Q.ABS_SPECIES(1).SX        = ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, 0.5, ...
                                          'lin', 0.2, 0.00, @log10 ) + ...
                      covmat1d_from_cfun( Q.ABS_SPECIES(1).GRIDS{1}, 0.1, ...
                                          'lin', 0.5, 0.00, @log10 );
  Q.ABS_SPECIES(1).MINMAX  = 1e-6;
end



%- Water
%
% This generates no absorption, as linefile has no H2O lines
```

```
%
Q.ABS_SPECIES(2).TAG       = { 'H2O' };
Q.ABS_SPECIES(2).RETRIEVE = false;
Q.ABS_SPECIES(2).ATMDATA  = gf_artsxml( fullfile( fascod, ...
     'midlatitude-winter', 'midlatitude-winter.H2O.xml' ), 'H2O', 'vmr_field' );




%- Wind
%
% Here demonstrated for v-component, that should the component of main
% concern for ground-based instruments. This component can be seen as the
% part of the horizontal wind going along the viewing direction, where
% positive values mean a movement away from the sensor.
%
Q.WIND_V_FIELD    = [];      % This is a short-cut for zero wind.
%
Q.WIND_V.RETRIEVE = false;   % Set to true to activate retrieval
Q.WIND_V.L2       = true;
Q.WIND_V.GRIDS    = { Q.P_GRID(1:2:end), [], [] };
Q.WIND_V.SX       = covmat1d_from_cfun( Q.WIND_V(1).GRIDS{1}, 40, ...
                                          'lin', 0.5, 0.00, @log10 );
if Q.WIND_V.RETRIEVE
  Q.WSMS_AT_START{end+1} = 'IndexSet(abs_f_interp_order,3)';
end




%- Pointing
%
% Here just included for testing purposes, of little interest for ground-based
% spectrometers.
%
Q.POINTING.RETRIEVE       = false;
Q.POINTING.DZA            = 0.1;
Q.POINTING.POLY_ORDER     = 0;
Q.POINTING.CALCMODE       = 'recalc';
Q.POINTING.SX             = 1;
Q.POINTING.L2             = true;




%- Frequency shift
%
Q.FSHIFTFIT.RETRIEVE      = true;
Q.FSHIFTFIT.DF            = 25e3;
Q.FSHIFTFIT.SX            = 50e3^2;
Q.FSHIFTFIT.L2            = true;


%- Frequency stretch
```

```
%
Q.FSTRETCHFIT.RETRIEVE      = false;   % Set to true to activate retrieval
Q.FSTRETCHFIT.DF            = 25e3;
Q.FSTRETCHFIT.SX            = 50e3^2;
Q.FSTRETCHFIT.L2            = true;


%- Polyfit
%
% A polynomial of order 3 is used for "baseline fit".
%
Q.POLYFIT.RETRIEVE          = true;
Q.POLYFIT.ORDER             = 3;
Q.POLYFIT.L2                = true;
Q.POLYFIT.SX0               = 1^2;
Q.POLYFIT.SX1               = 0.5^2;
Q.POLYFIT.SX2               = 0.2^2;
Q.POLYFIT.SX3               = 0.1^2;


%- Sinefit
%
% Here demonstrated with two period lengths
%
Q.SINEFIT.RETRIEVE          = false;   % Set to true to activate retrieval
Q.SINEFIT.PERIODS           = [ 75e3 200e3 ]';
Q.SINEFIT.L2                = true;
Q.SINEFIT.SX1               = 0.2^2;
Q.SINEFIT.SX2               = 0.4^2;

return
%-------------------------------------------------------------------------------
%-------------------------------------------------------------------------------
%-------------------------------------------------------------------------------




function Y = y_demo(Q)

% The data should be loaded from one or several files, but are here genereted
% by a forward model call to show how qpack2 can also be used to generate
% simulaled measurements (matching a priori assumptions).

% The simulated data model airborn measurements at two different zenith
% angles, from two nearby positions.

% Init Y
```

```
%
Y = qp2_y;


% Set a date
%
Y.YEAR  = 2008;
Y.MONTH = 2;
Y.DAY   = 25;


% Lat / lon
%
Y.LATITUDE  = 45;
Y.LONGITUDE = exp(1);


% An airborn measurement assumed here
%
Y.Z_PLATFORM = 10.5e3;
Y.ZA         = 50;


% Reference point for hydrostatic equilibrium
%
Y.HSE_P = 100e2;
Y.HSE_Z = 16e3;


% Set backend frequencies
%
Y.F = Q.SENSOR_RESPONSE.F_BACKEND;


% Thermal noise standard deviation
%
Y.TNOISE = 0.1;
% To test varying noise
%Y.TNOISE = linspace( 0.03, 0.07, length(Y.F) )';


% Simulate a measurement
%
Y.Y = [];                          % A flag to tell qpack2 to calculate the
%                                    spectrum ({} signifies undefined!).


% Add a second measurement
%
Y(2) = Y(1);
%
Y(2).LONGITUDE = pi;
Y(2).ZA        = 45;



% Calculate simulated spectra
%
```

```
Y = qpack2( Q, oem, Y );                % Dummy oem structure OK here


% Add thermal noise
%
% The correlation specified in Q is included
%
for i = 1 : length(Y)
  Y(i).Y = Y(i).Y + Y(i).TNOISE .* make_noise(1,Q.TNOISE_C);
end

% Add a constant "baseline shift" for measurement 2
%
Y(2).Y = Y(2).Y + 1;
```

B   The fields of Q

The information below is obtained in Atmlab by typing:

```
>> qinfo( @qarts )
```

C   The fields of O

The information below is obtained in Atmlab by typing:

```
>> qinfo( @oem )
```

D   The fields of Y

The information below is obtained in Atmlab by typing:

```
>> qinfo( @qp2_y )
```